

1 Multigrid for Poisson: $A\underline{u} = \underline{f}$.

- Multigrid methods are some of the fastest available solution strategies for multidimensional PDEs.
- In practice, they are most commonly used as preconditioners within Krylov-subspace projection (KSP) methods, such as conjugate gradient iteration or GMRES, which gives them speed and robustness.
- Within a KSP, the multigrid solver returns \underline{z} in the preconditioning step, *Solve*

$$M_g \underline{z} = \underline{r}. \tag{1}$$

- In the following, we'll think about solving $A\underline{u} = \underline{f}$ as being equivalent to $M_g \underline{z} = \underline{r}$, with recognition that our initial guess will be $\underline{u}_0 = 0$ and that \underline{u} (which approximates the error in the context of a KSP) will satisfy homogeneous boundary conditions.
- *There are two essential components to a multigrid solver,*
 - 1. Error Smoothing**, and
 - 2. Coarse-Grid Correction**
- We'll start with smoothing.

2 Smoothing

- Consider the following fixed-point iteration for $A\underline{u} = \underline{f}$, with $\underline{u}_0 = 0$:

for $k = 1 : m$

$$\underline{u}_k = \underline{u}_{k-1} + \omega M^{-1}(\underline{f} - A\underline{u}_{k-1}) \quad (2)$$

end

- Here, M is alternately referred to as a *smoother* or *preconditioner*.
- Most frequently, $M = \text{diag}(A)$.
- We also sometimes refer to the whole loop as a smoother.
- And, in the KSP, we refer to the whole multigrid part as a preconditioner (of the KSP).
- Let $\underline{r}_k := \underline{f} - A\underline{u}_k \equiv A(\underline{u} - \underline{u}_k) = A\underline{e}_k$.
- The *residual*, \underline{r}_k , is A times the *error*, $\underline{e}_k := \underline{u} - \underline{u}_k$. Always.
- Solving $A\underline{u} = \underline{b}$ amounts to finding the error.

- Suppose we have a known guess (or *iterate*), \underline{u}_k .
- Then $\underline{e}_k := \underline{u} - \underline{u}_k \iff \underline{u} = \underline{u}_k + \underline{e}_k$.
- If we find \underline{e}_k , we're done.
- The equation for \underline{e}_k is: $A\underline{e}_k = \underline{r}_k := \underbrace{\underline{f} - A\underline{u}_k}_{\text{computable}} \tag{3}$
- It is easy to compute the rhs, but solving $A\underline{e}_k = \underline{r}_k$ seems as difficult as solving $A\underline{u} = \underline{b}$.
- We shall see, however, that we can infer certain properties about \underline{e}_k that will make it relatively easy to solve.
- Specifically, we will design the fixed point iteration to try to make \underline{e}_k as *smooth* as possible.
- At that point, we can approximate it on a coarser mesh, which will mean that it is less expensive to compute.

- Let's look more closely at the smoothing step,

$$\underline{u}_k = \underline{u}_{k-1} + \omega M^{-1}(\underline{f} - A\underline{u}_{k-1})$$

- For now, think of $M = I$ or, almost equivalently, as a modified system with $\tilde{A} = M^{-1}A$ and $\tilde{f} = M^{-1}\underline{f}$, such that

$$\underline{u}_k = \underline{u}_{k-1} + \omega(\tilde{f} - \tilde{A}\underline{u}_{k-1}). \quad (4)$$

- We'll drop the \sim for now, but will re-introduce M later.
- The iteration (4) is known as *Richardson iteration*.
- We'll see that it is analogous to Euler-Forward timestepping.
- Recall the unsteady heat equation,

$$\frac{\partial u}{\partial t} = \nabla^2 u + f \implies \frac{du}{dt} = -Au + \underline{f}, \quad (5)$$

for which the unsteady part decays in time, with the particular property that the high wavenumber components decay rapidly, while the low wavenumber components decay more slowly.

- For small enough timestep sizes, Euler-Forward timestepping

$$\frac{\underline{u}^n - \underline{u}^{n-1}}{\Delta t} = -A\underline{u}^{n-1} + \underline{f} \quad (6)$$

$$\implies \underline{u}^n = \underline{u}^{n-1} + \Delta t(\underline{f} - A\underline{u}^{n-1}), \quad (7)$$

yields solutions with this same property.

- To see this behavior, let's discuss how timestepping is relevant to solving $A\underline{u} = \underline{f}$.
- Define \underline{u} as the solution to $A\underline{u} = \underline{f}$.
- We find an equation for the *error*, $\underline{e}^n := \underline{u} - \underline{u}^n$, as follows.

$$\begin{aligned} &+ (\underline{u} = \underline{u} + \Delta t(\underline{f} - A\underline{u})) \\ &- (\underline{u}^n = \underline{u}^{n-1} + \Delta t(\underline{f} - A\underline{u}^{n-1})) \\ \hline &\underline{e}^n = \underline{e}^{n-1} - \Delta t A \underline{e}^{n-1}. \end{aligned}$$

- If we start with $\underline{u}^n = 0$, this equation is nothing other than EF applied to the homogeneous ODE for the error,

$$\frac{d\underline{e}}{dt} = -A\underline{e}, \quad \underline{e}(t=0) = \underline{u}. \quad (8)$$

- The error decays from its initial value, $\underline{e}^0 = \underline{u}$ towards 0.

We've seen this characteristic behavior in earlier discussions of timestepping.




Figure 1: Euler-Forward solution behavior for unsteady 1D heat equation at two spatial resolutions, same Δt .

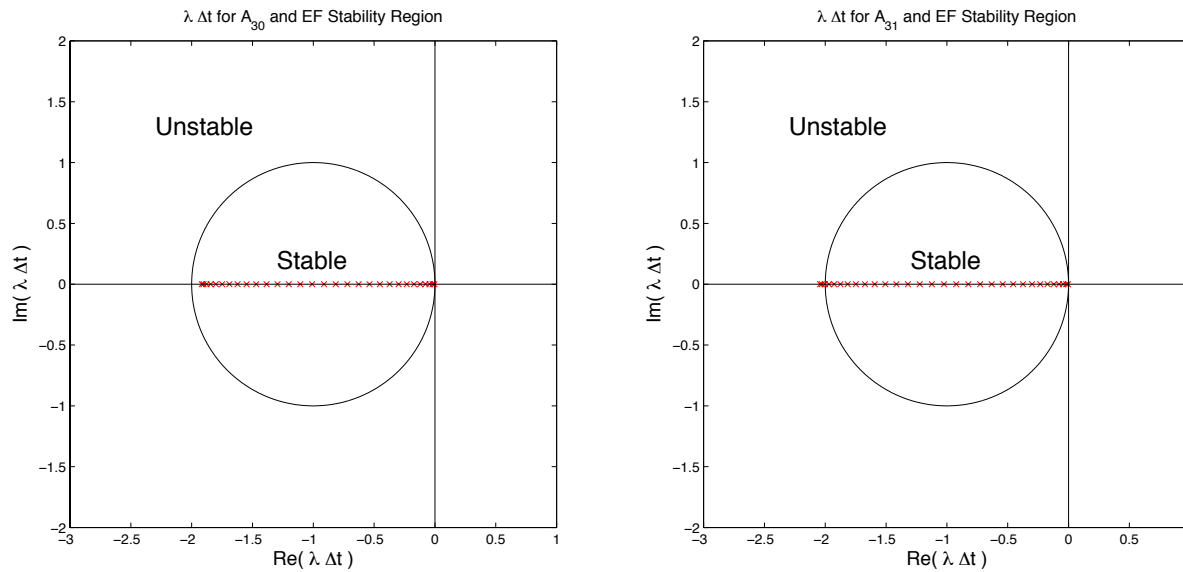


Figure 2: Stability region for Euler forward timestepping and the scaled eigenvalues, $\Delta t \lambda$ for $L = -\nu A$ with $\nu = .05$ and A the m -point 2nd-order finite difference approximation to $-u_{xx}$: (left) $m=30$, (right) $m=31$.

- In the figure on the upper left, the nonsmooth IC rapidly evolves to a smooth one, provided that $|\lambda\Delta t| < 2$, which is the bound for the EF stability region.
- If we exceed this bound, however, the solution blows up, as illustrated in the top right.
- So, the prefactor Δt in the EF timestepping depends on $\max \lambda(A)$,

$$\underline{u}^n = \underline{u}^{n-1} + \Delta t(\underline{f} - A\underline{u}^{n-1})$$

$$\Delta t \leq \frac{2}{\lambda_m},$$

assuming $0 < \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_m$.

- Notice that the remark regarding smoothing applies only to the spectral components that are well within the stability region.

- Recall our earlier analysis for EF with $L = -A$ and A an $m \times m$ SPD matrix having a complete set of eigenvectors.
- For any $\underline{u} \in \mathbb{R}^m$, we can write

$$\underline{u} = \sum_k \hat{u}_k \underline{z}_k, \quad (9)$$

$$L \underline{z}_k = \mu_k \underline{z}_k, \quad (10)$$

where $\mu_k(L) = -\lambda_k(A)$ are the eigenvalues of L

- Applying this decomposition to our EF timestepper, we have



$$\hat{u}_k^n = \underbrace{[1 + \Delta t \mu_k]}_{g(\mu \Delta t)} \hat{u}_k^{n-1}, \quad (11)$$

where $g = g_{EF}(\mu \Delta t)$ is the *growth factor* for EF.

- We have a similar growth factor for our exact (analytical) timestepper.
- If we do not discretize in time, the Fourier coefficients satisfy

$$\hat{u}_k^n = \underbrace{e^{\mu \Delta t}}_{\tilde{g}(\mu \Delta t)} \hat{u}_k^{n-1}. \quad (12)$$

- Consider the case when $\mu(L) = -\lambda(A)$ is negative real (corresponding to $\lambda(A) > 0$).
- We plot g_{EF} and \tilde{g} below for two different distributions of $\mu\Delta t$.



- In the right figure, we show a distribution of $\mu_k\Delta t$ in red where the largest (in magnitude) is $\mu_m\Delta t = -2$.
- For this case, the growth rate will be $g_{EF}(-2) = -1$, which means that high-frequency error will *not decay*.

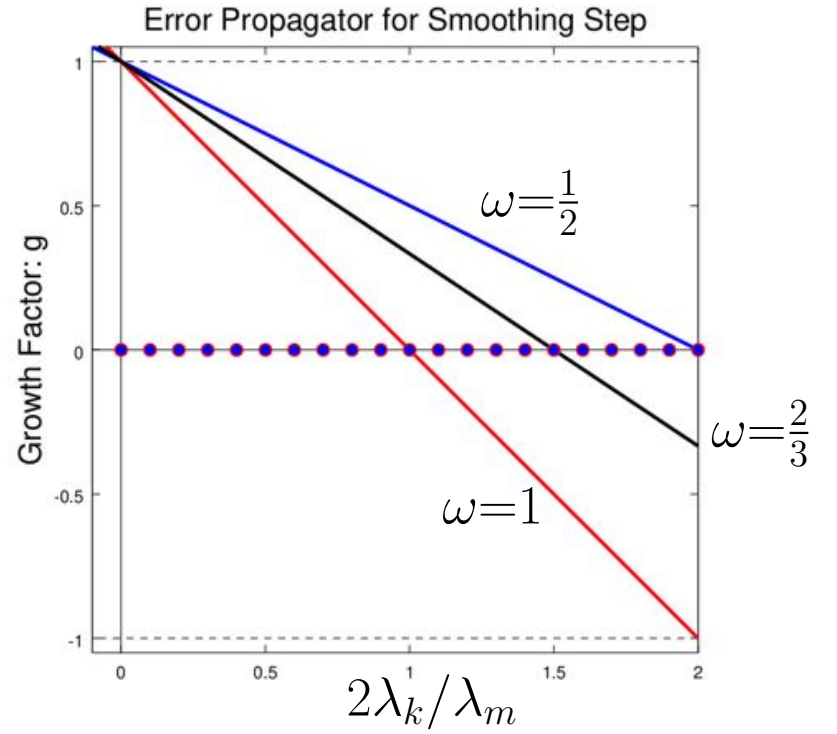
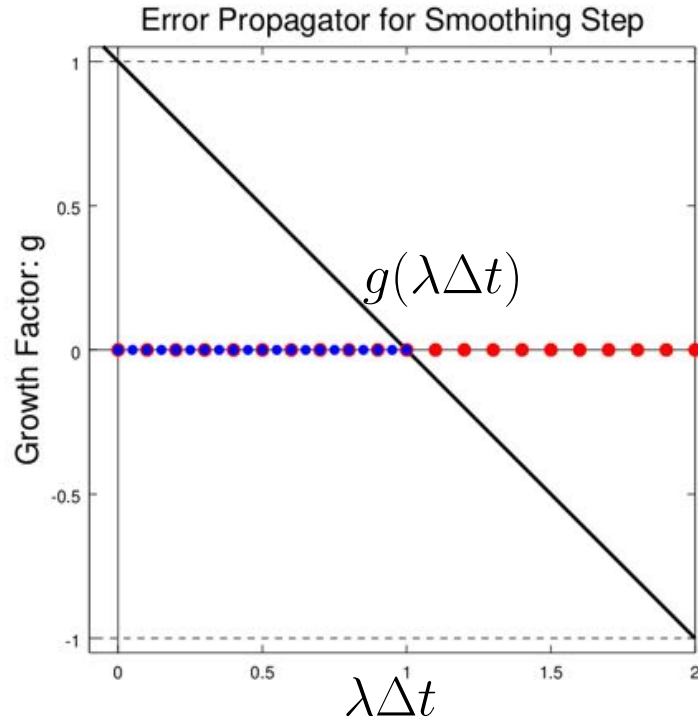
- In the figure on the right, we illustrate a case with the same eigenvalue distribution, μ_k , but with a value of Δt such that $\mu_m \Delta t = -1$.
- For this case, $g_{EF}(\mu_m \Delta t) = 0$, which means that the high-frequency error is annihilated after a single timestep.
- Paradoxically, taking a *smaller* timestep leads to more rapid decay of the high-frequency error.
- We refer to this process as under-relaxed Jacobi (or Richardson) iteration.

3 From Euler Forward to Smoothing

- Let's turn the preceding plots around so that the x -axis is $\lambda_k(A)$ and focus solely on the EF case.
- The growth factor is then

$$g(\lambda \Delta t) = 1 - \lambda \Delta t, \tag{13}$$

which we plot on left below on the left for the same two eigenvalue distributions considered in the preceding figures.



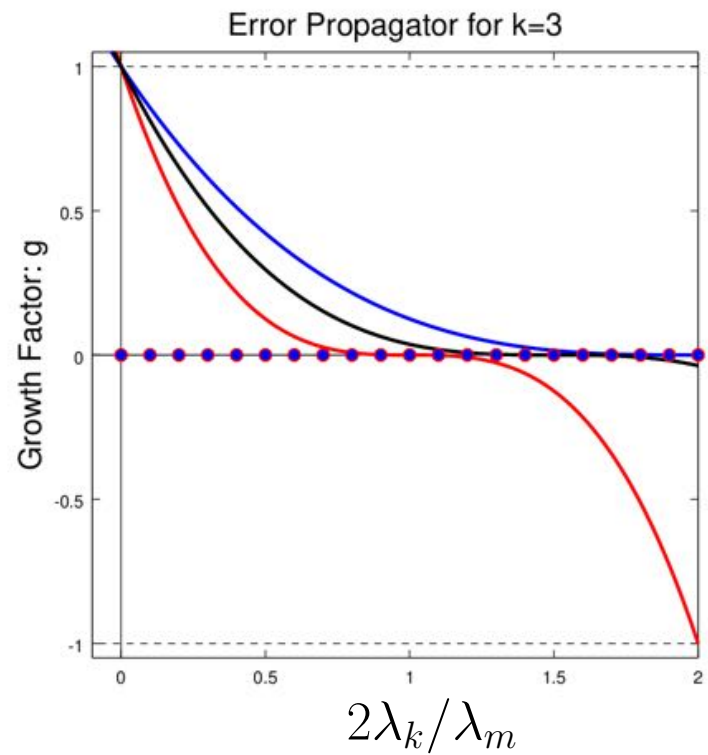
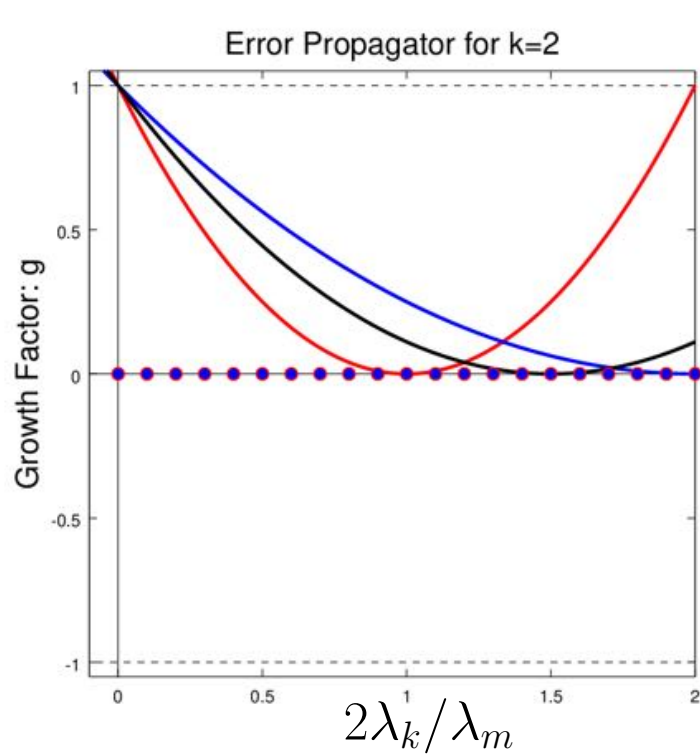
- In the right figure, we have rescaled the x -axis to be $2\lambda_k/\lambda_m$.
- The growth factor in this case becomes

$$g_\omega(\lambda_k) = 1 - \omega \frac{2\lambda_k}{\lambda_m}. \quad (14)$$

- The eigenvalue distributions on the x -axis no longer depend on Δt , but the growth curves do, with

$$\Delta t = \omega \frac{2}{\lambda_m}. \quad (15)$$

- Here, $\omega = 1$ corresponds to the largest stable Δt (i.e., the *nonsmoothing case* corresponding to the red dots in the left figure)
- $\omega = 1/2$ corresponds to $\Delta t = 1/\lambda_m$, which rapidly annihilates the high wavenumber error (i.e., *smooths* the solution).
- A third case, $\omega = \frac{2}{3}$ is generally more optimal than $\omega = \frac{1}{2}$.
- With this choice, each round of the smoothing iteration reduces the error in the upper half of the spectrum by $1/3$.
- The figures below illustrate what happens to the error distribution after two and three rounds of smoothing.



- It is clear from the black curve that $\omega = \frac{2}{3}$ does a better job of suppressing the error in the mid-range of the spectrum, albeit at the cost of less error reduction than $\omega = \frac{1}{2}$ at the very high end.

- Let's look at the smoothing idea in the context of the SEM.
- Here, λ reflects the eigenvalues of $D^{-1}A$, where $D = \text{diag}(A)$.
- The `smoother.m` file reads,


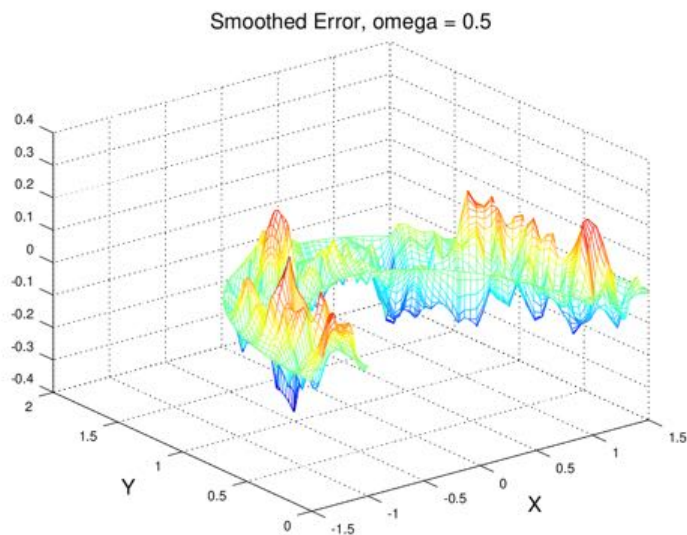
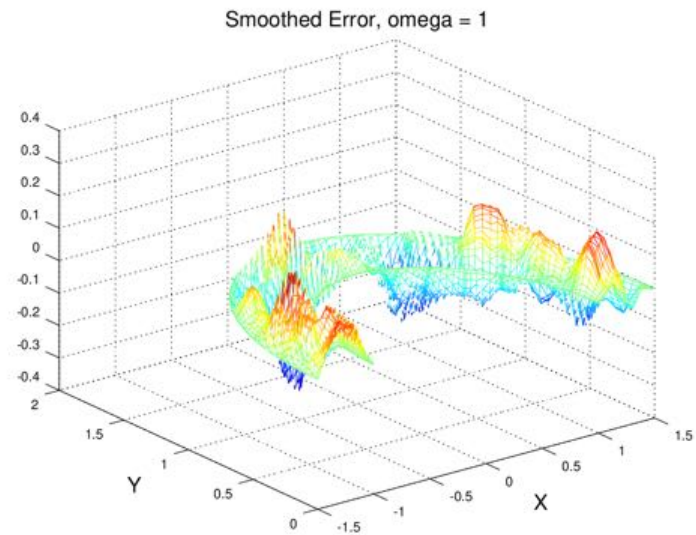
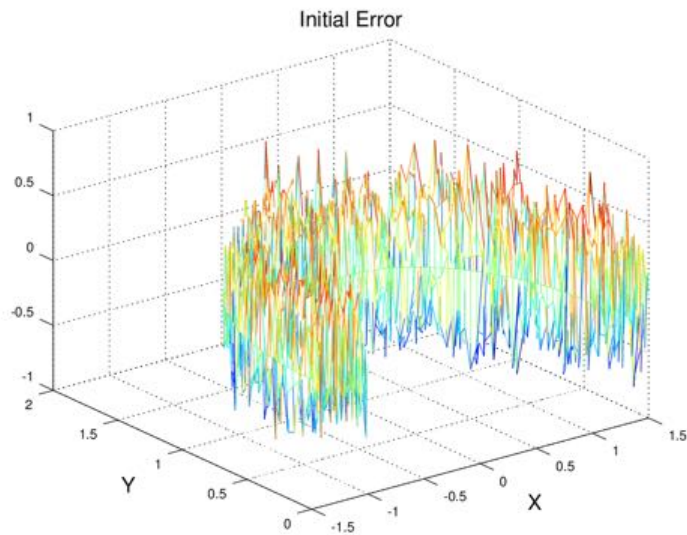
```
function [u]=smoother...
    (F1,lam_max,omega,m,b0,nu,M,Q,B1,Grr,Grs,Gss,Dh,Di,ifnull);

omega = omega * (2 / lam_max); %% Rescale omega

r = F1;
z = Di.*(M.*qqt(Q,r));    %% diagonal preconditioner
u = omega*z;              %% initial guess for x is 0.

for iter=2:m;              %% Apply m rounds of Jacobi smoothing
    r = F1 - ax1(u,b0,nu,B1,Grr,Grs,Gss,Dh);
    z = Di.*(M.*qqt(Q,r));
    u = u + omega*z;
end;
```



- Here, \underline{u} is the current guess,
 $\underline{r} := \underline{f} - A\underline{u}$ is the residual, and
 $\underline{z} = M^{-1}\underline{r}$ is the search direction.
- The update step is thus $\underline{u} = \underline{u} + \omega M^{-1}(\underline{f} - A\underline{u})$.
- The following figures show the results of $k = 50$ rounds of smoothing for $\omega = 1, \frac{1}{2}$, and $\frac{2}{3}$, along with the initial error, for a 2D SEM example.



- The solution (initial error), upper left, is a random field with the correct boundary conditions.
- The upper right figure clearly shows that the choice $\omega=1$ fails to smooth the solution.
- The lower left figure shows that $\omega=\frac{1}{2}$ leads to an excellent smoother.
- However, the case $\omega=\frac{2}{3}$, lower right, leads to more overall error reduction.

4 Chebyshev Smoothing

- A significant advance in multigrid methods has resulted from recognizing that smoothing can be more effective if one chooses different values of ω on each iteration.
- The idea is to spread out the zero crossings such that the error is suppressed over a broader range, rather than just focusing on error reduction in the neighborhood of the principal root.
- Below, we plot the $k = 1$ plots for our three choices of ω on the left.
- On the right we contrast smoothing after 3 rounds with our optimized $\omega = \frac{2}{3}$ (red) and the case with $\omega_k = \frac{1}{2}, \frac{2}{3}$, and 1, for round $k=1, 2$, and 3 (blue).



- Overall, the blue curve shows broader error suppression than the red curve.
- Another distribution, seen as the black dashed line, shows better suppression using roots $[0.9, 1.4, 1.9]$.
- This observation leads to the natural question, *Can we further optimize the choice of the roots of our k th-order polynomial?*
- Remarkably, the answer has its roots in spectral methods!
- Namely, if you wish to minimize the maximum of a polynomial over a given interval, subject to some nontrivial scaling, then answer is typically a scaled and translated *Chebyshev* polynomial.
- There are a variety of choices, but recent work by James Lottes has developed an algorithm using *optimized 4th-kind Chebyshev polynomials*, which is given below.

```

function [x]=cheby4(Fl,lam_max,omega,k,b0,nu,M,Q,B1,Grr,Grs,Gss,Dh,Di,ifnull);
%% Apply kth-order optimized 4th-kind Chebyshev
%% (from Malachi Phillips dissertation,2023).
lmax_i = 1./lam_max;
beta=Beta(k,:);
r = Fl;
x = 0*r;
d = (lmax_i*4./3.)*(Di.*(M.*qqt(Q,r)));%Diagonal preconditioner
for i=1:k-1;
    x = x + beta(i)*d;
    r = r - axl(d,b0,nu,B1,Grr,Grs,Gss,Dh);
    s1= (2*i-1)/(2*i+3); s2=lmax_i*(8*i+4)/(2*i+3);
    d = s1*d + s2*Di.*(M.*qqt(Q,r));
end;
x = x + beta(k)*d;

```

- The algorithm requires an estimate of $\lambda_m(M^{-1}A)$, which is readily obtained from power iteration or from running CG in “Lanczos” mode, which can give estimates of the extreme eigenpairs.
- It also requires some precomputed values of β_i , which are provided in functional and tabular form in recent articles.
- Here is the table used in our matlab code:

```

Beta=[ ...
1.125000000000000      0      0      0      0      0      0;
1.02387287570313 1.26408905371085      0      0      0      0      0;
1.00842544782028 1.08867839208730 1.33753125909618      0      0      0      0;
1.00391310427285 1.04035811188593 1.14863498546254 1.38268869241000      0      0      0;
1.00212930146164 1.02173711549260 1.07872433192603 1.19810065292663 1.41322542791682      0      0;
1.00128517255940 1.01304293035233 1.04678215124113 1.11616489419675 1.23829020218444 1.43524297106744      0;
1.00083464397912 1.00843949430122 1.03008707768713 1.07408384092003 1.15036186707366 1.27116474046139 1.45186658649364];

```

- We’ll look at the impact of this smoother momentarily.

5 Coarse-Grid Correction

- Once the *error* is smooth, we can return to our modified problem, *Solve* $A\underline{e}_k = \underline{r}_k$.
- Given that this problem arises in the context of an outer iteration, we do not need to solve it exactly.
- Since $e_k(\mathbf{x})$ is a smooth function, we can represent the solution on a coarser grid.
- In the SEM context with polynomial order N , we might consider approximating \underline{e}_k (actually, $e_k(\mathbf{x})$) with polynomial order $N_c = N/2$ or $N - 2$, say.
- The natural approach is to use a Galerkin approximation.
- Let $\tilde{\underline{e}}_c := J\underline{e}_c$ be the coarse-grid approximation to \underline{e}_k .
- Here, J is an interpolation matrix from the coarse space (order $N_c < N$) to the fine space (order N).
- The variational problem is, *Find* $\tilde{\underline{e}}_c \in \mathcal{R}(J)$ *such that, for all* $\underline{v} \in \mathcal{R}(J)$,

$$\underline{v}^T A \tilde{\underline{e}}_c = \underline{v}^T \underline{r}_k \equiv \underline{v}^T A \underline{e}_k. \quad (16)$$

- With $\underline{v} = J\underline{v}_c$ and $\tilde{\underline{e}}_c := J\underline{e}_c$, this becomes *Find* $\underline{e}_c \in \mathbb{R}^{n_c}$ *such that, for all* $\underline{v}_c \in \mathbb{R}^{n_c}$,

$$\underline{v}_c^T J^T A J \underline{e}_c = \underline{v}_c^T J^T \underline{r}_k. \quad (17)$$

or simply, with $A_c := J^T A J$,

$$A_c \underline{e}_c = J^T \underline{r}_k =: \underline{r}_c. \quad (18)$$

- It is easy to interpolate on an element-by-element basis, so application of J and J^T is *completely local*.
- Unfortunately, $A_c = J^T A J$ loses the tensor-product structure of A .
- We therefore replace it with A_{N_c} , the system matrix associated with the lower polynomial degree, N_c .
- Once we solve $A_c \underline{e}_c = \underline{r}_c$, we interpolate to the fine mesh and add this correction to the current (smoothed) iterate:

Coarse-Grid Correction: $\underline{u}_k = \underline{u}_k + J \underline{e}_c, \quad \underline{e}_c := A_c^{-1} J^T \underline{r}_k.$

- Depending on the size, we solve the coarse grid problem either iteratively or directly.
- The most common approach is to apply multigrid to this problem as well, which changes the *two-level multigrid* algorithm we've described so far into an actual *multigrid* method.
- In the following, examples we revisit our SEM matlab example using two-level multigrid with $N = 7$, $N_c = 5$, and either 5 rounds of damped Jacobi smoothing $\omega = \frac{2}{3}$ or optimized 4th-kind Chebyshev smoothing with $k = 6$.
- Note that the work for $k = 6$ is effectively the same as 5 rounds of damped Jacobi smoothing.

Damped Jacobi, $\omega=2/3$			Optimized 4th-kind Chebyshev		
	e_smooth	e_coarse		e_smooth	e_coarse
1	6.3368e-01	2.5709e-01	1	5.7602e-01	2.6874e-01
2	2.0752e-01	1.7723e-01	2	1.9069e-01	1.0420e-01
3	1.5486e-01	1.2586e-01	3	9.0264e-02	4.2108e-02
4	1.1029e-01	1.0912e-01	4	3.9747e-02	1.9548e-02
5	9.3898e-02	8.0521e-02	5	1.8402e-02	1.1508e-02
6	6.8317e-02	7.0165e-02	6	1.0927e-02	5.5164e-03
7	5.9945e-02	5.3564e-02	7	5.2486e-03	3.3857e-03
8	4.5350e-02	4.5779e-02	8	3.2528e-03	1.6877e-03
9	3.8958e-02	3.5762e-02	9	1.6222e-03	1.0741e-03
10	3.0253e-02	3.0143e-02	10	1.0269e-03	5.4384e-04

- For each smoother choice, the table shows the error after the smoothing step and after the coarse grid correction for each MG step.
- We can see that, for the same work, the 4th-kind Chebyshev significantly outperforms, *by $3\times$* , standard Jacobi smoothing.
- Although not discussed here, the 4th-kind Chebyshev is more robust than red-black Gauss-Seidel smoothing for geometries that have high aspect-ratio cells.
(Also, Gauss-Seidel is not amenable to matrix-free formulations.)
- Importantly, this preconditioner will be wrapped with a KSP, which will significantly accelerate its performance.

- Another observation: For this example, setting a tighter coarse-solve tolerance significantly improves the Chebyshev case, but not the Jacobi case.
- This observation leads to the idea that one can adjust the tolerance for the coarse (PCG) solver to be a fraction of the initial coarse-PCG residual, as illustrated by the use of **nrc**, the 2-norm of the coarse residual—normalized by the domain volume, in the following code snippet.

```
%% Coarse-grid correction

r1 = r1-ax1(Us,b0,nu,B1,Grr,GrS,Gss,Dh);      %% Update residual
rc = tensor3(Jc',1,Jc',r1);
nrc= sqrt( sum(sum(sum(M.*dA.*(r1.^2))))/vol );
tolc=0.1*nrc;
[Ec,iterc,res,lam_crs]=...
    pcg_lambda(rc,tolc,100,b0,nu,Mc,Qc,B1c,Grrc,Grsc,Gssc,Dc,dAc,ifnull);
E1 = tensor3(Jc,1,Jc,Ec);

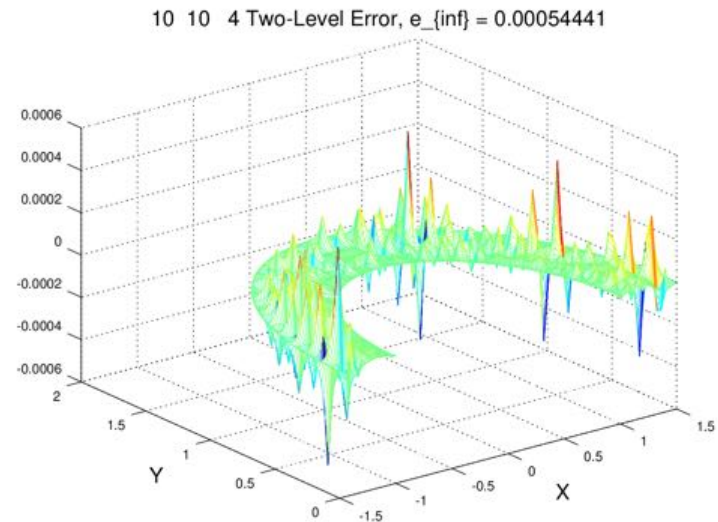
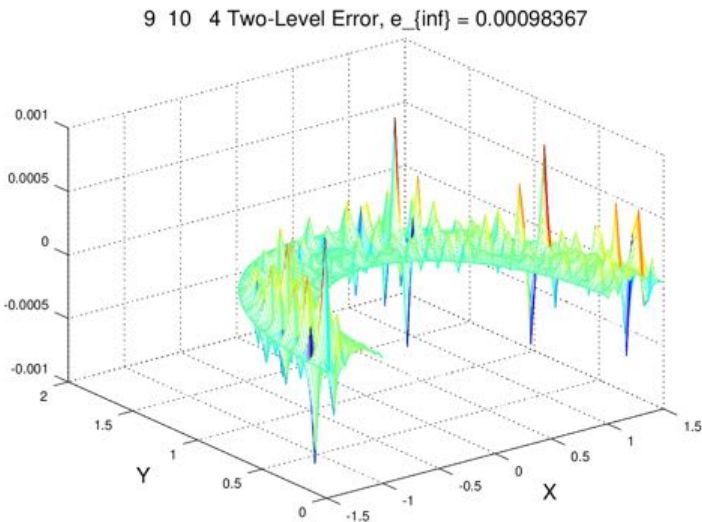


Umg = Umg + E1;                                %% Add coarse-grid correction
```

- The output for this strategy is shown below, where we see the coarse iteration counts, **kc** and the rapid convergence of Chebyshev-accelerated p-multigrid (pMG). (pMG = coarsening by reducing local polynomial order, $p = N$.)
- The driver, **slide_table.m**, and supporting scripts are posted to the [Relate](#) page. (Note that the exact solution is *random*, so that tables are a bit different whenever this demo is run.)

Damped Jacobi, $\omega=2/3$						
N	Nc	k	e_inf(k)	smoother	tol_c	kc
10	4	1	0.35707	0.87416	5.58806	18
10	4	2	0.27208	0.31388	1.57768	9
10	4	3	0.18824	0.25511	0.83471	13
10	4	4	0.15520	0.15888	0.54382	3
10	4	5	0.13856	0.14039	0.39838	2
10	4	6	0.12369	0.12547	0.31063	2
10	4	7	0.10579	0.11221	0.25203	6
10	4	8	0.09492	0.09576	0.21026	2
10	4	9	0.08524	0.08605	0.17690	2
10	4	10	0.07666	0.07742	0.15070	2

Optimized 4th-kind Chebyshev						
N	Nc	k	e_inf(k)	smoother	tol_c	kc
10	4	1	0.70037	0.75592	8.17176	9
10	4	2	0.12425	0.57576	1.30327	17
10	4	3	0.03994	0.09235	0.25363	15
10	4	4	0.02072	0.02760	0.06939	19
10	4	5	0.01275	0.01350	0.02363	6
10	4	6	0.00766	0.00810	0.01022	12
10	4	7	0.00461	0.00495	0.00514	6
10	4	8	0.00263	0.00298	0.00284	15
10	4	9	0.00156	0.00168	0.00162	6
10	4	10	0.00094	0.00100	0.00095	11

- In the tables, **smoother** refers to the error after the smoothing step, whereas **e_inf(k)** refers to the error after the k th pMG sweep.



- The figures above show the errors for a Chebyshev pMG case with $N = 10$ and $N_c = 4$. The initial error is in the top left and the error for iterations $k = 8, 9$, and 10 in the other panels.
- We see that the error decreases for $k = 8$ to 10 , but the similarity of the error distributions indicates that there is significant potential to accelerate the process by projecting the solution at each iteration to avoid resolving the same problem.